

Logistic Regression and Softmax Regression

Prof. Mingkui Tan

SCUT Machine Intelligence Laboratory (SMIL)



Contents

- 1 Logistic Regression
- 2 Softmax Regression
- 3 Variant of Softmax Loss

Contents

1 Logistic Regression

2 Softmax Regression

3 Variant of Softmax Loss

Data Example

Dataset: $\mathcal{D} = \{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)\}$

■ $\mathbf{x}_i \leftarrow$ health information

■ $y_i = \pm 1 \leftarrow$ did he have a heart attack or not

■ Given the health information of one person:


age	62 years
gender	male
blood sugar	120 mg/dL40,000
HDL	50
LDL	120
Mass	190 lbs
Height	5' 10''
...	...

How to infer the **probability of heart attack?**


Linear Classification and Regression

The linear signal:

$$z = \mathbf{w}^T \mathbf{x}$$



Linear Classification



Linear Regression

Probability Function

- To infer the probability of heart attack $P[y = +1|\mathbf{x}]$, the **probability function** of logistic function is as follows:


$$h_{\mathbf{w}}(\mathbf{x}) = g(z) = g\left(\sum_{i=1}^m w_i x_i\right) = g(\mathbf{w}^T \mathbf{x})$$

Here, $z = \mathbf{w}^T \mathbf{x}$, $g(\cdot)$ is a **logistic function**:

$$g(z) = \frac{1}{1 + e^{-z}}$$

Properties of Logistic Function

$$g(z) = \frac{1}{1 + e^{-z}}$$



- The function is a continuous function
- If $z \rightarrow +\infty$, then $g(z) \rightarrow 1$; if $z \rightarrow -\infty$, then $g(z) \rightarrow 0$

$$g(z) = \frac{1}{1 + e^{-z}} = \frac{e^z}{1 + e^z}$$

$$g(-z) = \frac{1}{1 + e^z} = 1 - g(z)$$

How to Learn **w**?

- Intuitively, similar to SVM, we need to define a **Loss Function** to find a good $h_{\mathbf{w}}(\mathbf{x})$ so that it fits the following targets well:

$$h_{\mathbf{w}}(\mathbf{x}) \text{ is good if : } \begin{cases} h_{\mathbf{w}}(\mathbf{x}) \approx 1, & y = 1 \\ h_{\mathbf{w}}(\mathbf{x}) \approx 0, & y = -1 \end{cases}$$

- Can we use the least square loss below?

$$\mathcal{L}(\mathbf{w}) = \frac{1}{n} \sum_{i=1}^n (h_{\mathbf{w}}(\mathbf{x}_i) - \frac{1}{2}(1 + y_i))^2$$

Questions: Why the least square loss is in this way?

- Can we use this loss? The answer is **Negative!** Why?
- Probabilily $h_{\mathbf{w}}(\mathbf{x})=1.001$ which is better than $h_{\mathbf{w}}(\mathbf{x})=0.9$
- But $h_{\mathbf{w}}(\mathbf{x})$ denotes the probability, thus $h_{\mathbf{w}}(\mathbf{x})$ must satisfy:

$$h_{\mathbf{w}}(\mathbf{x}) \leq 1.$$

How to Learn **w**?

- We need to define a **Loss Function** to find a good $h_{\mathbf{w}}(\mathbf{x})$ so that it fits the following targets well:

$$h_{\mathbf{w}}(\mathbf{x}) \text{ is good if : } \begin{cases} h_{\mathbf{w}}(\mathbf{x}) \approx 1, & y = 1 \\ h_{\mathbf{w}}(\mathbf{x}) \approx 0, & y = -1 \end{cases}$$

- The least square loss is **no longer valid** here since $h_{\mathbf{w}}(\mathbf{x})$ is a **probability function** with $h_{\mathbf{w}}(\mathbf{x}) \leq 1$.
- Here, we introduce a **new loss** called **logistic loss** as below:

$$\mathcal{L}(\mathbf{w}) = \frac{1}{n} \sum_{i=1}^n \log(1 + e^{-y_i \mathbf{w}^T \mathbf{x}_i})$$

Why the logistic loss is in this form?

Probabilistic View of Training Samples

- Recall $h_{\mathbf{w}}(\mathbf{x})$ is a **probability function** to predict the probability of an instance \mathbf{x} being to the label $y_i \in \{-1, 1\}$ as below:

$$P(y|\mathbf{x}) = \begin{cases} g(\mathbf{w}^T \mathbf{x}), & y = 1 \\ 1 - g(\mathbf{w}^T \mathbf{x}) = g(-\mathbf{w}^T \mathbf{x}), & y = -1 \end{cases}$$

- The training sample (\mathbf{x}_i, y_i) can be considered as **random variables** sampled from a sample space $\{\mathcal{X}, \mathcal{Y}\}$.
- The instance \mathbf{x}_i and its label y_i follow a **conditional probability**:

$$P(y_i|\mathbf{x}_i) = g(y_i \mathbf{w}^T \mathbf{x}_i)$$

The label y_i is definitely determined by the observation \mathbf{x}_i , namely **y_i is condition on \mathbf{x}_i**

How to Learn **w**?

Recall that the training samples $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)$ can be considered as random variables following the a **conditional probability** as below:

$$P(y_i|\mathbf{x}_i) = g(y_i \mathbf{w}^T \mathbf{x}_i)$$

Likelihood of training examples:

Assume that $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)$ are **independently** sampled, the joint distribution (or likelihood) $P(y_1, \dots, y_n | \mathbf{x}_1, \dots, \mathbf{x}_n)$ of $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)$ satisfies:

$$P(y_1, \dots, y_n | \mathbf{x}_1, \dots, \mathbf{x}_n) = \prod_{i=1}^n P(y_i | \mathbf{x}_i)$$

How to Learn \mathbf{w} ?

Note the parameter \mathbf{w} determines the distribution

$$P(y_i|\mathbf{x}_i) = g(y_i\mathbf{w}^T\mathbf{x}_i)$$

- Given the likelihood $P(y_1, \dots, y_n|\mathbf{x}_1, \dots, \mathbf{x}_n) = \prod_{i=1}^n P(y_i|\mathbf{x}_i)$, we can estimate \mathbf{w} with **Maximum Likelihood Estimation (MLE)**
- What is **Maximum Likelihood Estimation**?

Definition: Maximum Likelihood Estimation

Maximum Likelihood Estimation (MLE) is a statistical method used to make inferences about parameters of the underlying probability distribution of a given data set.

How to estimate parameter \mathbf{w} in $h_{\mathbf{w}}(\mathbf{x})$ with MLE?

How to Learn **w**?

Estimate **w** by maximizing the likelihood

$$\max_{\mathbf{w}} P(y_1, \dots, y_n | \mathbf{x}_1, \dots, \mathbf{x}_n) = \prod_{i=1}^n P(y_i | \mathbf{x}_i)$$

$$\max \prod_{i=1}^n P(y_i | \mathbf{x}_i) \Leftrightarrow \max \log \left(\prod_{i=1}^n P(y_i | \mathbf{x}_i) \right)$$

$$\equiv \max \sum_{i=1}^n \log P(y_i | \mathbf{x}_i)$$

$$\Leftrightarrow \min -\frac{1}{n} \sum_{i=1}^n \log P(y_i | \mathbf{x}_i)$$

How to Learn \mathbf{w} ?

Estimate \mathbf{w} by **maximizing the likelihood** $\max_{\mathbf{w}} \prod_{i=1}^n P(y_i | \mathbf{x}_i)$

$$\begin{aligned} \max_{\mathbf{w}} \prod_{i=1}^n P(y_i | \mathbf{x}_i) &\Leftrightarrow \max \log \left(\prod_{i=1}^n P(y_i | \mathbf{x}_i) \right) \\ &\Leftrightarrow \min -\frac{1}{n} \sum_{i=1}^n \log P(y_i | \mathbf{x}_i) \\ &\Leftrightarrow \min \frac{1}{n} \sum_{i=1}^n \log \frac{1}{P(y_i | \mathbf{x}_i)} \quad \leftarrow P(y_i | \mathbf{x}_i) = g(y_i \mathbf{w}^T \mathbf{x}_i) \\ &\equiv \min \frac{1}{n} \sum_{i=1}^n \log \frac{1}{g(y_i \mathbf{w}^T \mathbf{x}_i)} \quad \leftarrow g(z) = \frac{1}{1 + e^{-z}} \\ &\equiv \min \frac{1}{n} \sum_{i=1}^n \log(1 + e^{-y_i \mathbf{w}^T \mathbf{x}_i}) \equiv \min \mathcal{L}(\mathbf{w}) \end{aligned}$$

Definition: Logistic regression

$$\max_{\mathbf{w}} \prod_{i=1}^n P(y_i | \mathbf{x}_i) = \min_{\mathbf{w}} \mathcal{L}(\mathbf{w}) = \min_{\mathbf{w}} \frac{1}{n} \sum_{i=1}^n \log(1 + e^{-y_i \cdot \mathbf{w}^T \mathbf{x}_i})$$

Regularization Required

Similar to SVM, we employ **Regularization** to avoid overfitting issue

- We have the following objective function for logistic regression:

$$J(\mathbf{w}) = \mathcal{L}(\mathbf{w}) + \frac{\lambda}{2} ||\mathbf{w}||_2^2$$

Here, $\mathcal{L}(\mathbf{w}) = \frac{1}{n} \sum_{i=1}^n \log(1 + e^{-y_i \cdot \mathbf{w}^T \mathbf{x}_i})$ is called **Logistic Loss** and λ is the regularization parameter.

Why need regularization?

- “Simple” model
- Less prone to overfitting

SVM vs Logistic Regression

- SVM:

$$\min J(\mathbf{w}) = \frac{1}{n} \sum_{i=1}^n \max(0, 1 - y_i \mathbf{w}^T \mathbf{x}_i) + \frac{\lambda}{2} \|\mathbf{w}\|_2^2$$

- logistic regression:

$$\min J(\mathbf{w}) = \frac{1}{n} \sum_{i=1}^n \log(1 + e^{-y_i \mathbf{w}^T \mathbf{x}_i}) + \frac{\lambda}{2} \|\mathbf{w}\|_2^2$$

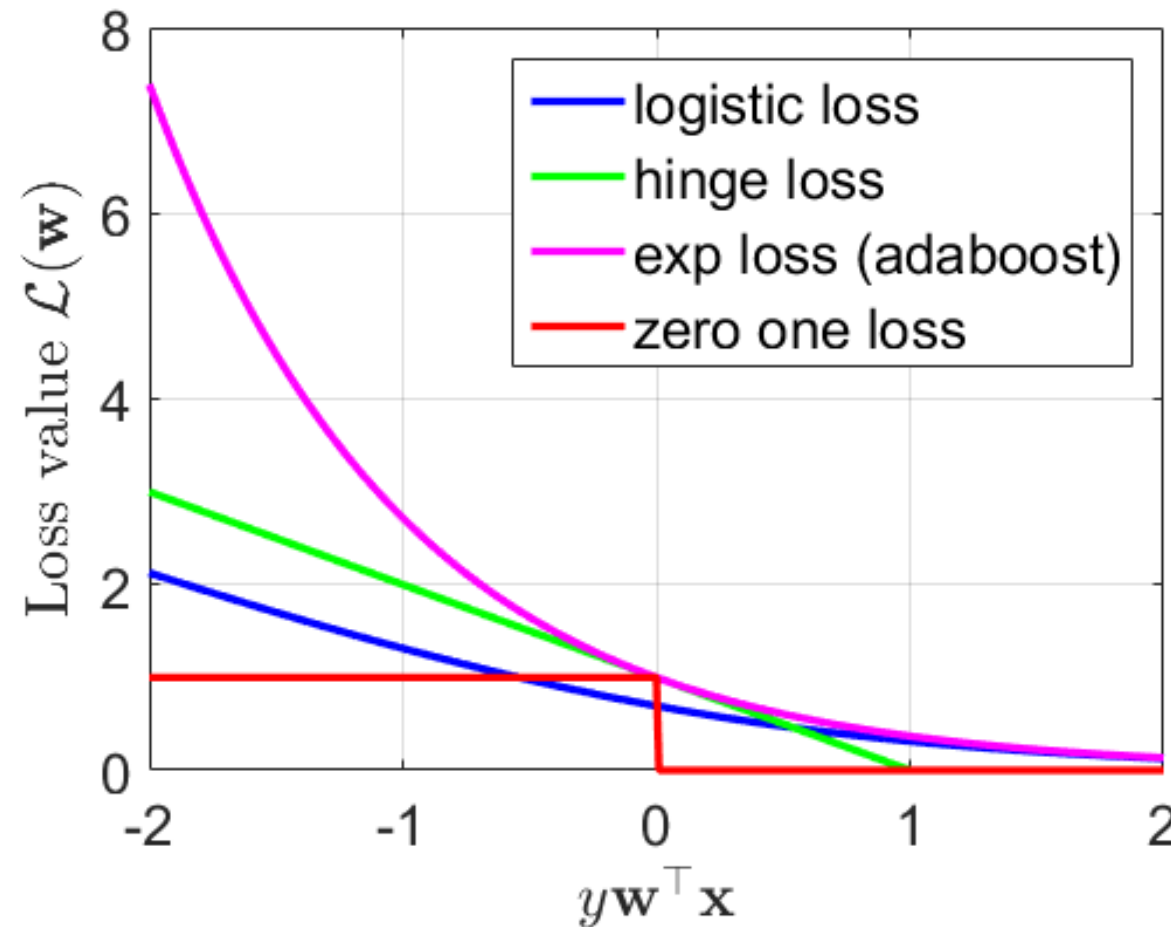
- The regularization term $\|\mathbf{w}\|_2^2$ is called L_2^2 regularizer

- Connections to SVM:

- Both are supervised algorithms

- Both are used to solve **binary classification** problem

Graphical Comparison of Loss Functions



Comparison of Different Loss Functions

logistic loss:

$$\mathcal{L}(\mathbf{x}_i; \mathbf{w}) = \log(1 + e^{-y_i \mathbf{w}^T \mathbf{x}_i})$$

hinge loss:

$$\mathcal{L}(\mathbf{x}_i; \mathbf{w}) = \max(0, 1 - y_i \mathbf{w}^T \mathbf{x}_i)$$

exponential loss
(for adaboost):



$$\mathcal{L}(\mathbf{x}_i; \mathbf{w}) = e^{-y_i \mathbf{w}^T \mathbf{x}_i}$$

zero one loss:

$$\mathcal{L}(\mathbf{x}_i; \mathbf{w}) = \begin{cases} 0, & y_i \mathbf{w}^T \mathbf{x}_i > 0 \\ 1, & y_i \mathbf{w}^T \mathbf{x}_i \leq 0 \end{cases}$$

Graphical Comparison of Three Loss Functions

- 绿色: 0/1 损失
- 蓝色: SVM Hinge 损失函数
- 红色: Logistic 损失函数



How to Learn \mathbf{w} ?


Minimize $J(\mathbf{w})$ by (Stochastic) Gradient Descent: $\min_{\mathbf{w}} J(\mathbf{w})$

- Compute gradient $\frac{\partial J(\mathbf{w})}{\partial \mathbf{w}}$ of $J(\mathbf{w})$ with respect to \mathbf{w} :

$$\frac{\partial J(\mathbf{w})}{\partial \mathbf{w}} = -\frac{1}{n} \sum_{i=1}^n \frac{y_i \mathbf{x}_i e^{-y_i \mathbf{w}^T \mathbf{x}_i}}{1 + e^{-y_i \mathbf{w}^T \mathbf{x}_i}} + \lambda \mathbf{w}$$

- Update parameters **with learning rate** η

$$\mathbf{w} := \mathbf{w} - \eta \frac{\partial J(\mathbf{w})}{\partial \mathbf{w}}$$



Note:
$$\begin{aligned} \frac{\partial J(\mathbf{w})}{\partial \mathbf{w}} &= \frac{1}{n} \sum_{i=1}^n \frac{\partial(\log(1 + e^{-y_i \mathbf{w}^T \mathbf{x}_i}))}{\partial \mathbf{w}} + \lambda \mathbf{w} = \frac{1}{n} \sum_{i=1}^n \frac{1}{1 + e^{-y_i \mathbf{w}^T \mathbf{x}_i}} \cdot \frac{\partial(e^{-y_i \mathbf{w}^T \mathbf{x}_i})}{\partial \mathbf{w}} + \lambda \mathbf{w} \\ &= \frac{1}{n} \sum_{i=1}^n \frac{1}{1 + e^{-y_i \mathbf{w}^T \mathbf{x}_i}} \cdot e^{-y_i \mathbf{w}^T \mathbf{x}_i} \cdot (-y_i \mathbf{x}_i) + \lambda \mathbf{w} = -\frac{1}{n} \sum_{i=1}^n \frac{y_i \mathbf{x}_i e^{-y_i \mathbf{w}^T \mathbf{x}_i}}{1 + e^{-y_i \mathbf{w}^T \mathbf{x}_i}} + \lambda \mathbf{w} \end{aligned}$$

Logistic Regression for $y_i \in \{0,1\}$

Previous study considers $y_i \in \{-1, +1\}$, but what if $y_i \in \{0,1\}$ and what if $y_i \in \{0, 1, \dots, K - 1\}$?

- Let us first consider the simple case: $y_i \in \{0,1\}$
- Similar to the case $y_i \in \{-1,1\}$, we define the probability of \mathbf{x}_i being with the label $y_i \in \{0,1\}$ as follows:

$$P(y_i|\mathbf{x}_i) = \begin{cases} h_{\mathbf{w}}(\mathbf{x}_i), & y = 1 \\ 1 - h_{\mathbf{w}}(\mathbf{x}_i), & y = 0 \end{cases}$$

$$\text{Where } h_{\mathbf{w}}(\mathbf{x}) = g(\mathbf{w}^T \mathbf{x}) = \frac{1}{1 + e^{-\mathbf{w}^T \mathbf{x}}}$$

- More specifically, the instance \mathbf{x}_i and its label y_i follow the conditional probability as below:

$$P(y_i|\mathbf{x}_i) = h_{\mathbf{w}}(\mathbf{x}_i)^{y_i} \cdot (1 - h_{\mathbf{w}}(\mathbf{x}_i))^{1-y_i}$$

Again, Resort to Maximum Likelihood Estimation

Note the parameter \mathbf{w} determines the distribution

$$P(y_i|\mathbf{x}_i) = h_{\mathbf{w}}(\mathbf{x}_i)^{y_i} \cdot (1 - h_{\mathbf{w}}(\mathbf{x}_i))^{1-y_i}$$

Likelihood of training examples:

Assuming that $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)$ are **independently** sampled, the joint distribution (or likelihood) $P(y_1, \dots, y_n | \mathbf{x}_1, \dots, \mathbf{x}_n)$ of $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)$ satisfies $P(y_1, \dots, y_n | \mathbf{x}_1, \dots, \mathbf{x}_n) = \prod_{i=1}^n P(y_i | \mathbf{x}_i)$

■ We can estimate \mathbf{w} with **Maximum Likelihood Estimation (MLE)**

Similar to $y_i \in \{-1, 1\}$, we **maximize the likelihood** to estimate \mathbf{w}

$$\max_{\mathbf{w}} P(y_1, \dots, y_n | \mathbf{x}_1, \dots, \mathbf{x}_n) = \prod_{i=1}^n P(y_i | \mathbf{x}_i)$$

How to Learn **w**?

Similar to $y_i \in \{-1, 1\}$, we maximize the likelihood to estimate \mathbf{w}

$$\max_{\mathbf{w}} P(y_1, \dots, y_n | \mathbf{x}_1, \dots, \mathbf{x}_n) = \prod_{i=1}^n P(y_i | \mathbf{x}_i)$$

$$\begin{aligned} \max \prod_{i=1}^n P(y_i | \mathbf{x}_i) &\Leftrightarrow \max \log \left(\prod_{i=1}^n P(y_i | \mathbf{x}_i) \right) \\ &\equiv \max \sum_{i=1}^n \log P(y_i | \mathbf{x}_i) \\ &\Leftrightarrow \min -\frac{1}{n} \sum_{i=1}^n \log P(y_i | \mathbf{x}_i) \end{aligned}$$

How to Learn **w**?

Estimate **w** by **maximizing the likelihood** $\max_{\mathbf{w}} \prod_{i=1}^n P(y_i | \mathbf{x}_i)$

$$\begin{aligned} \max \prod_{i=1}^n P(y_i | \mathbf{x}_i) &\Leftrightarrow \max \log \left(\prod_{i=1}^n P(y_i | \mathbf{x}_i) \right) \\ &\equiv \min -\frac{1}{n} \sum_{i=1}^n \log \left(h_{\mathbf{w}}(\mathbf{x}_i)^{y_i} \cdot (1 - h_{\mathbf{w}}(\mathbf{x}_i))^{1-y_i} \right) \\ &\equiv \min -\frac{1}{n} \sum_{i=1}^n (y_i \log h_{\mathbf{w}}(\mathbf{x}_i) + (1 - y_i) \log(1 - h_{\mathbf{w}}(\mathbf{x}_i))) \\ &\equiv \min \mathcal{L}(\mathbf{w}) \end{aligned}$$

Regularization Required

We employ **Regularization** to avoid overfitting issue

- We have the following objective function for logistic regression:

$$J(\mathbf{w}) = \mathcal{L}(\mathbf{w}) + \frac{\lambda}{2} \|\mathbf{w}\|_2^2$$

Now, the **logistic loss** becomes

$$\mathcal{L}(\mathbf{w}) = -\frac{1}{n} \sum_{i=1}^n (y_i \log h_{\mathbf{w}}(\mathbf{x}_i) + (1 - y_i) \log(1 - h_{\mathbf{w}}(\mathbf{x}_i)))$$

Minimize $J(\mathbf{w})$ by (Stochastic) Gradient Descent: $\min_{\mathbf{w}} J(\mathbf{w})$

How to Learn \mathbf{w} ?


Minimize $J(\mathbf{w})$ by (Stochastic) Gradient Descent: $\min_{\mathbf{w}} J(\mathbf{w})$

- Compute gradient $\frac{\partial J(\mathbf{w})}{\partial \mathbf{w}}$ of $J(\mathbf{w})$ with respect to \mathbf{w} :

$$\frac{\partial J(\mathbf{w})}{\partial \mathbf{w}} = \frac{1}{n} \sum_{i=1}^n (h_{\mathbf{w}}(\mathbf{x}_i) - y_i) \mathbf{x}_i + \lambda \mathbf{w}$$

- Update parameters with **learning rate** η

$$\mathbf{w} := \mathbf{w} - \eta \frac{\partial J(\mathbf{w})}{\partial \mathbf{w}}$$



Details of Calculate $\frac{\partial J(\mathbf{w})}{\partial \mathbf{w}}$

Note:

$$\begin{aligned}\frac{\partial J(\mathbf{w})}{\partial \mathbf{w}} &= \frac{1}{n} \sum_{i=1}^n \left(-y_i \cdot \frac{1}{h_{\mathbf{w}}(\mathbf{x}_i)} \cdot \frac{\partial h_{\mathbf{w}}(\mathbf{x}_i)}{\partial \mathbf{w}} + (1 - y_i) \cdot \frac{1}{1 - h_{\mathbf{w}}(\mathbf{x}_i)} \frac{\partial h_{\mathbf{w}}(\mathbf{x}_i)}{\partial \mathbf{w}} \right) + \lambda \mathbf{w} \\ &= \frac{1}{n} \sum_{i=1}^n \left(-y_i \cdot \frac{1}{h_{\mathbf{w}}(\mathbf{x}_i)} \cdot \frac{\partial h_{\mathbf{w}}(\mathbf{x}_i)}{\partial \mathbf{w}} + (1 - y_i) \cdot \frac{1}{1 - h_{\mathbf{w}}(\mathbf{x}_i)} \frac{\partial h_{\mathbf{w}}(\mathbf{x}_i)}{\partial \mathbf{w}} \right) + \lambda \mathbf{w} \\ &= \frac{1}{n} \sum_{i=1}^n \left(-y_i \cdot \frac{\mathbf{x}_i h_{\mathbf{w}}(\mathbf{x}_i) (1 - h_{\mathbf{w}}(\mathbf{x}_i))}{h_{\mathbf{w}}(\mathbf{x}_i)} + (1 - y_i) \cdot \frac{\mathbf{x}_i h_{\mathbf{w}}(\mathbf{x}_i) (1 - h_{\mathbf{w}}(\mathbf{x}_i))}{1 - h_{\mathbf{w}}(\mathbf{x}_i)} \right) + \lambda \mathbf{w} \\ &= \frac{1}{n} \sum_{i=1}^n (h_{\mathbf{w}}(\mathbf{x}_i) - y_i) \mathbf{x}_i + \lambda \mathbf{w}\end{aligned}$$

Contents

1 Logistic Regression

2 Softmax Regression

3 Variant of Softmax Loss

Extension to Multi-class Classification

Previous study considers $y \in \{0,1\}$, but what if $y \in \{0,1, \dots, K-1\}$?

Dataset: $\mathcal{D} = \{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)\}$


■ \mathbf{x}_i is the observation for the i^{th} instance

■ $y_i \in \{0,1, \dots, K-1\}$ is the label for the i^{th} instance

■ Task: Predict the probability of a testing instance \mathbf{x} being to any class $j \in \{0,1, \dots, K-1\}$ as o_j

■ Then o_j must follow:

$$0 \leq o_j \leq 1, \quad \sum_j o_j = 1$$



Softmax Regression for Multi-class Classification

To handle **multi-class** task, for each class $j \in \{0, \dots, K - 1\}$, we define a weight vector \mathbf{w}_j associated with this class

■ $\mathbf{W} := [\mathbf{w}_0 \ \mathbf{w}_1 \ \dots \ \mathbf{w}_{K-1}]$ is a **matrix** of K weight **vectors**

$$\mathbf{W} = \begin{bmatrix} | & | & | & | \\ \mathbf{w}_0 & \mathbf{w}_1 & \cdots & \mathbf{w}_{K-1} \\ | & | & | & | \end{bmatrix}_{m \times K}$$

Here, m is the dimension of the sample, K is the number of classes

■ Let $z_j = \mathbf{w}_j^T \mathbf{x}$. We define the probability of an instance \mathbf{x} being to any class $j \in \{0, 1, \dots, K - 1\}$ as:

$$o_j = P(y = j | \mathbf{x}; \mathbf{W}) = \frac{e^{z_j}}{\sum_{l=0}^{K-1} e^{z_l}} = \frac{e^{\mathbf{w}_j^T \mathbf{x}}}{\sum_{l=0}^{K-1} e^{\mathbf{w}_l^T \mathbf{x}}}$$

Softmax Regression for Multi-class Classification


- Recall that the probability of an instance \mathbf{x} being to any class j is:

$$o_j = P(y = j|\mathbf{x}; \mathbf{W}) = \frac{e^{z_j}}{\sum_{l=0}^{K-1} e^{z_l}} = \frac{e^{\mathbf{w}_j^T \mathbf{x}}}{\sum_{l=0}^{K-1} e^{\mathbf{w}_l^T \mathbf{x}}}$$

- The function $\frac{e^{z_j}}{\sum_{l=0}^{K-1} e^{z_l}}$ is called **Softmax function**, where $\sum_{l=0}^{K-1} e^{z_l}$ is a normalization term to make all the elements **be summed to 1**

- Obviously, o_j follows:

$$0 \leq o_j \leq 1, \sum_j o_j = 1$$



Softmax Regression for Multi-class Classification

- For an instance \mathbf{x} , it can belong to any class j with probability:

$$H_{\mathbf{W}}(\mathbf{x}) = \begin{bmatrix} P(y = 0|\mathbf{x}; \mathbf{W}) \\ \vdots \\ P(y = j|\mathbf{x}; \mathbf{W}) \\ \vdots \\ P(y = K - 1|\mathbf{x}; \mathbf{W}) \end{bmatrix} = \frac{1}{\sum_{l=0}^{K-1} e^{\mathbf{w}_l^T \mathbf{x}}} \begin{bmatrix} e^{\mathbf{w}_0^T \mathbf{x}} \\ \vdots \\ e^{\mathbf{w}_j^T \mathbf{x}} \\ \vdots \\ e^{\mathbf{w}_{K-1}^T \mathbf{x}} \end{bmatrix}$$

- **Prediction:** Given any parameters \mathbf{W} , we can predict the label by:

$$\text{Prediction: } \hat{y} = \operatorname{argmax}_{j \in \{0, 1, \dots, K-1\}} P(y = j|\mathbf{x}; \mathbf{W})$$

How to learn a good \mathbf{W} to ensure correct prediction?

Cross-Entropy Loss for Multi-class Classification

- To learn $\mathbf{W} := [\mathbf{w}_0 \ \mathbf{w}_1 \ \dots \ \mathbf{w}_{K-1}]$, relying on the softmax function, we introduce the following **Cross-Entropy loss**:

$$\mathcal{L}(\mathbf{W}) = -\frac{1}{n} \left[\sum_{i=1}^n \sum_{j=0}^{K-1} \mathbb{I}\{y_i = j\} \log \frac{e^{\mathbf{w}_j^T \mathbf{x}_i}}{\sum_{l=0}^{K-1} e^{\mathbf{w}_l^T \mathbf{x}_i}} \right]$$

where $\mathbb{I}\{\cdot\}$ is the indicator function as follows:

$$\mathbb{I}\{A\} = \begin{cases} 1, & \text{if } A \text{ is a true statemet} \\ 0, & \text{if } A \text{ is a false statemet} \end{cases}$$

- The cross-entropy loss can be derived by Maximum Likelihood Estimation (MLE). Here, we omit the details.

Regularization Required

We employ **Regularization** to avoid overfitting issue

- We have the following **objective function** for softmax regression:

$$J(\mathbf{W}) = \mathcal{L}(\mathbf{W}) + \frac{\lambda}{2} ||\mathbf{W}||_2^2$$

Here, $\mathcal{L}(\mathbf{W}) = -\frac{1}{n} \left[\sum_{i=1}^n \sum_{j=0}^{K-1} \mathbb{I}\{y_i = j\} \log \frac{e^{\mathbf{w}_j^T \mathbf{x}_i}}{\sum_{l=0}^{K-1} e^{\mathbf{w}_l^T \mathbf{x}_i}} \right]$ is called

Cross-Entropy Loss and λ is the regularization parameter.

- Update parameters **W** by (Stochastic) Gradient Descent:

$$\mathbf{W} := \mathbf{W} - \eta \frac{\partial J(\mathbf{W})}{\partial \mathbf{W}}$$

How to compute $\frac{\partial J(\mathbf{W})}{\partial \mathbf{W}}$?

How to compute $\frac{\partial J(\mathbf{W})}{\partial \mathbf{W}}$?

- For \mathbf{w}_j ($j = 0, \dots, K - 1$), $\frac{\partial J(\mathbf{W})}{\partial \mathbf{w}_j}$ can be computed as follows:

$$\begin{aligned}\frac{\partial J(\mathbf{W})}{\partial \mathbf{w}_j} &= \frac{\partial \left\{ -\frac{1}{n} \left[\sum_{i=1}^n \sum_{j=0}^{K-1} \mathbb{I}\{y_i=j\} \log \frac{e^{\mathbf{w}_j^T \mathbf{x}_i}}{\sum_{l=0}^{K-1} e^{\mathbf{w}_l^T \mathbf{x}_i}} \right] + \frac{\lambda}{2} \|\mathbf{W}\|_2^2 \right\}}{\partial \mathbf{w}_j} \\&= -\frac{1}{n} \sum_{i=1}^n \frac{\partial \sum_{j=0}^{K-1} \mathbb{I}\{y_i=j\} (\log e^{\mathbf{w}_j^T \mathbf{x}_i} - \log \sum_{l=0}^{K-1} e^{\mathbf{w}_l^T \mathbf{x}_i})}{\partial \mathbf{w}_j} + \lambda \mathbf{w}_j \\&= -\frac{1}{n} \sum_{i=1}^n \left[\mathbb{I}\{y_i = j\} \mathbf{x}_i - \frac{1}{\sum_{l=0}^{K-1} e^{\mathbf{w}_l^T \mathbf{x}_i}} \cdot \frac{\partial \sum_{l=0}^{K-1} e^{\mathbf{w}_l^T \mathbf{x}_i}}{\partial \mathbf{w}_j} \right] + \lambda \mathbf{w}_j \\&= -\frac{1}{n} \sum_{i=1}^n \left[\mathbb{I}\{y_i = j\} \mathbf{x}_i - \frac{\mathbf{x}_i e^{\mathbf{w}_j^T \mathbf{x}_i}}{\sum_{l=0}^{K-1} e^{\mathbf{w}_l^T \mathbf{x}_i}} \right] + \lambda \mathbf{w}_j \\&= \frac{1}{n} \sum_{i=1}^n (P(y_i = j | \mathbf{x}_i; \mathbf{W}) - \mathbb{I}\{y_i = j\}) \mathbf{x}_i + \lambda \mathbf{w}_j\end{aligned}$$

Example of Softmax Regression

Softmax Classifier (Multinomial Logistic Regression)



Want to interpret raw classifier scores as **probabilities**


$$Z = f(x_i; W)$$

$$P(Y = k|X = x_i) = \frac{e^{s_k}}{\sum_j e^{s_j}} \quad \text{Softmax Function}$$

Probabilities
must be ≥ 0

Probabilities
must sum to 1

$$L_i = -\log P(Y = y_i|X = x_i)$$



Softmax Regression for Binary Classification

Previous cases consider softmax regression for multi-class classification. Can we use it for binary classification i.e., a special case where $K = 2$?

■ Recall that an instance \mathbf{x} can belong to any class j with probability:

$$H_{\mathbf{W}}(\mathbf{x}) = \begin{bmatrix} P(y = 0|\mathbf{x}; \mathbf{W}) \\ \vdots \\ P(y = j|\mathbf{x}; \mathbf{W}) \\ \vdots \\ P(y = K - 1|\mathbf{x}; \mathbf{W}) \end{bmatrix} = \frac{1}{\sum_{l=0}^{K-1} e^{\mathbf{w}_l^T \mathbf{x}}} \begin{bmatrix} e^{\mathbf{w}_0^T \mathbf{x}} \\ \vdots \\ e^{\mathbf{w}_j^T \mathbf{x}} \\ \vdots \\ e^{\mathbf{w}_{K-1}^T \mathbf{x}} \end{bmatrix}$$

■ When $K = 2$, we have:

$$H_{\mathbf{W}}(\mathbf{x}) = \begin{bmatrix} P(y = 0|\mathbf{x}; \mathbf{W}) \\ P(y = 1|\mathbf{x}; \mathbf{W}) \end{bmatrix} = \frac{1}{e^{\mathbf{w}_0^T \mathbf{x}} + e^{\mathbf{w}_1^T \mathbf{x}}} \begin{bmatrix} e^{\mathbf{w}_0^T \mathbf{x}} \\ e^{\mathbf{w}_1^T \mathbf{x}} \end{bmatrix}$$

Then, softmax regression is reduced to logistic regression

Softmax Regression for Binary Classification

- Recall that the weight matrix is $\mathbf{W} := [\mathbf{w}_0 \ \mathbf{w}_1]$
- When $K = 2$, we have

$$\begin{aligned} H_{\mathbf{W}}(\mathbf{x}) &= \begin{bmatrix} P(y = 0 | \mathbf{x}; \mathbf{W}) \\ P(y = 1 | \mathbf{x}; \mathbf{W}) \end{bmatrix} \\ &= \frac{1}{e^{\mathbf{w}_0^T \mathbf{x}} + e^{\mathbf{w}_1^T \mathbf{x}}} \begin{bmatrix} e^{\mathbf{w}_0^T \mathbf{x}} \\ e^{\mathbf{w}_1^T \mathbf{x}} \end{bmatrix} \\ &= \frac{1}{e^{(\mathbf{w}_0 - \mathbf{w}_1)^T \mathbf{x}} + e^{(\mathbf{w}_1 - \mathbf{w}_1)^T \mathbf{x}}} \begin{bmatrix} e^{(\mathbf{w}_0 - \mathbf{w}_1)^T \mathbf{x}} \\ e^{(\mathbf{w}_1 - \mathbf{w}_1)^T \mathbf{x}} \end{bmatrix} \\ &= \frac{1}{e^{(\mathbf{w}_0 - \mathbf{w}_1)^T \mathbf{x}} + e^{(0)^T \mathbf{x}}} \begin{bmatrix} e^{(\mathbf{w}_0 - \mathbf{w}_1)^T \mathbf{x}} \\ e^{(0)^T \mathbf{x}} \end{bmatrix} \end{aligned}$$

Softmax Regression for Binary Classification

■ Let $-\mathbf{w} = \mathbf{w}_0 - \mathbf{w}_1$, $H_{\mathbf{w}}(\mathbf{x}) = \begin{bmatrix} P(y = 0|\mathbf{x}; \mathbf{W}) \\ P(y = 1|\mathbf{x}; \mathbf{W}) \end{bmatrix}$

$$= \frac{1}{1 + e^{-\mathbf{w}^T \mathbf{x}}} \begin{bmatrix} e^{-\mathbf{w}^T \mathbf{x}} \\ 1 \end{bmatrix}$$

Probability in Logistic Regression:

$$P(y|\mathbf{x}) = \begin{cases} 1 - h_{\mathbf{w}}(\mathbf{x}), & y = 0 \\ h_{\mathbf{w}}(\mathbf{x}), & y = 1 \end{cases}$$
$$= \begin{bmatrix} 1 - \frac{1}{1 + e^{-\mathbf{w}^T \mathbf{x}}} \\ \frac{1}{1 + e^{-\mathbf{w}^T \mathbf{x}}} \end{bmatrix}$$
$$= \begin{bmatrix} 1 - h_{\mathbf{w}}(\mathbf{x}) \\ h_{\mathbf{w}}(\mathbf{x}) \end{bmatrix}$$

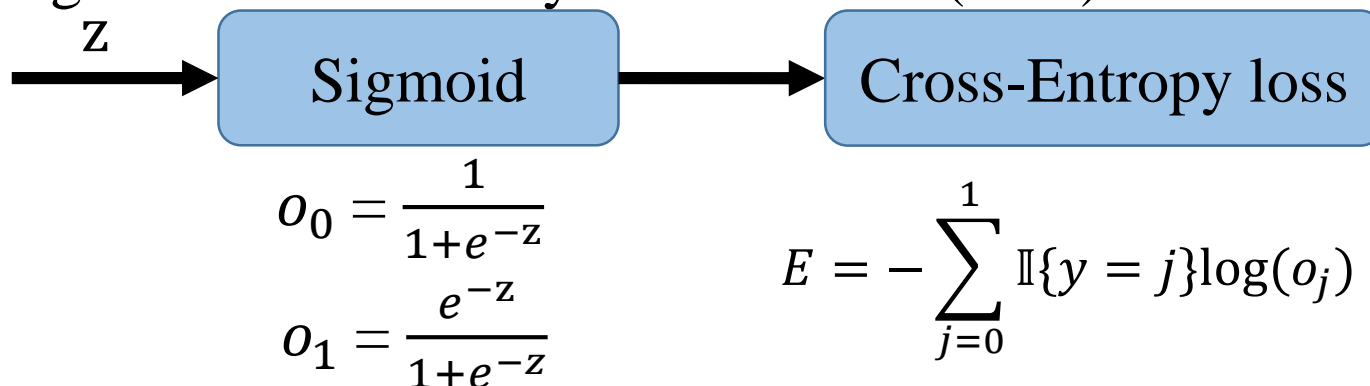
Logistic regression is a special case of softmax regression

Logistic Loss vs Softmax Cross-Entropy Loss

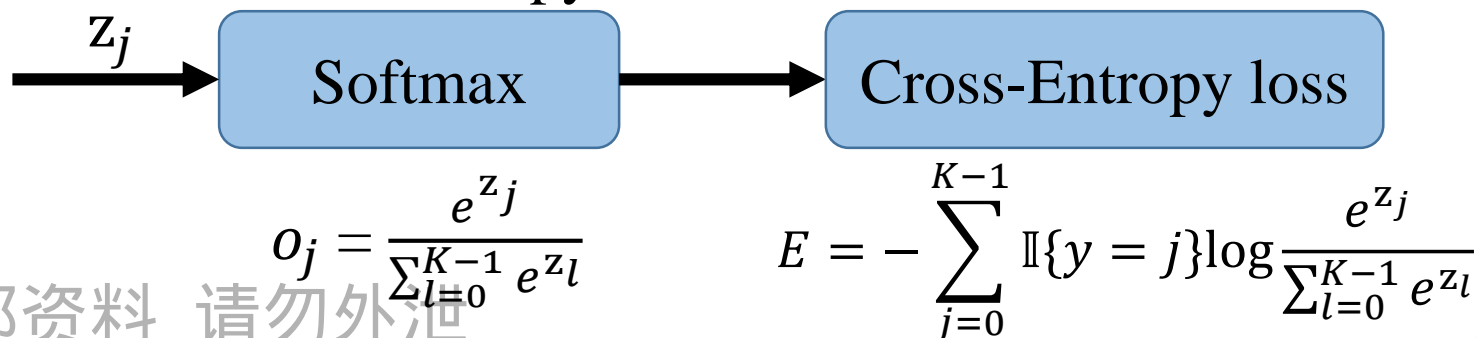
Cross-Entropy loss:

$$E = - \sum_{j=0}^{K-1} \mathbb{I}\{y = j\} \log(o_j)$$

■ Logistic loss for binary classification ($K=2$):



■ Softmax Cross-Entropy loss for multi-class classification:



Contents

1 Logistic Regression

2 Softmax Regression

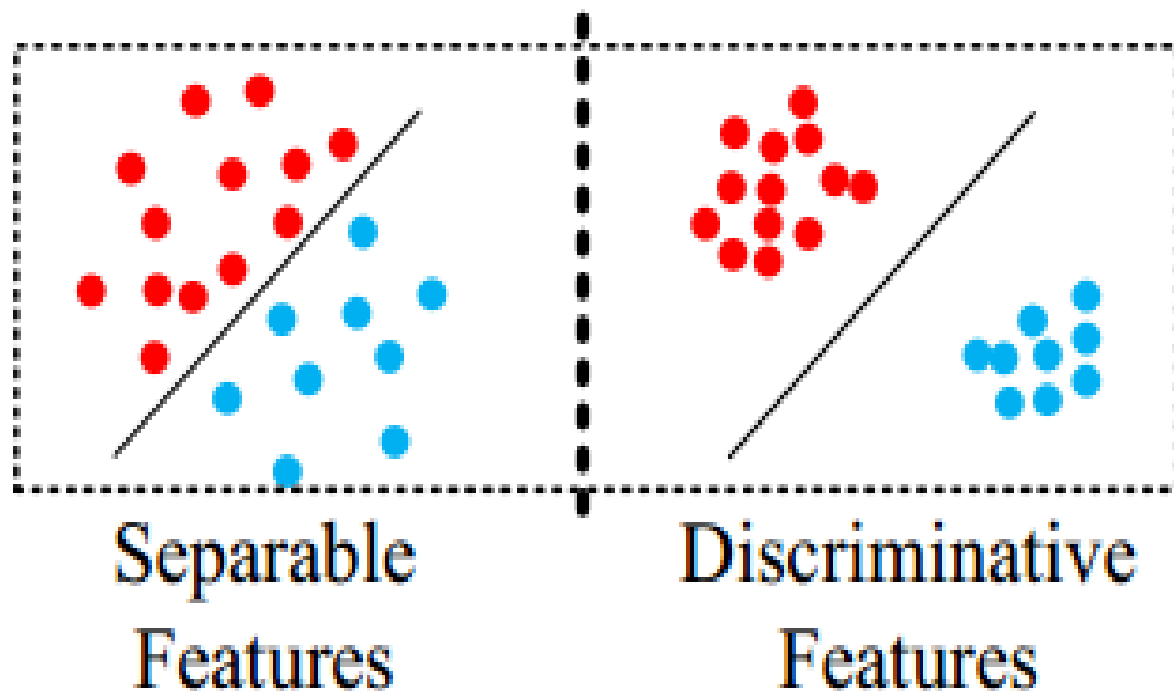
3 Variant of Softmax Loss

Two Variants of the Softmax Loss

- **Large-Margin Softmax Loss**
- Angular Softmax Loss

Motivation

- Learn discriminative features



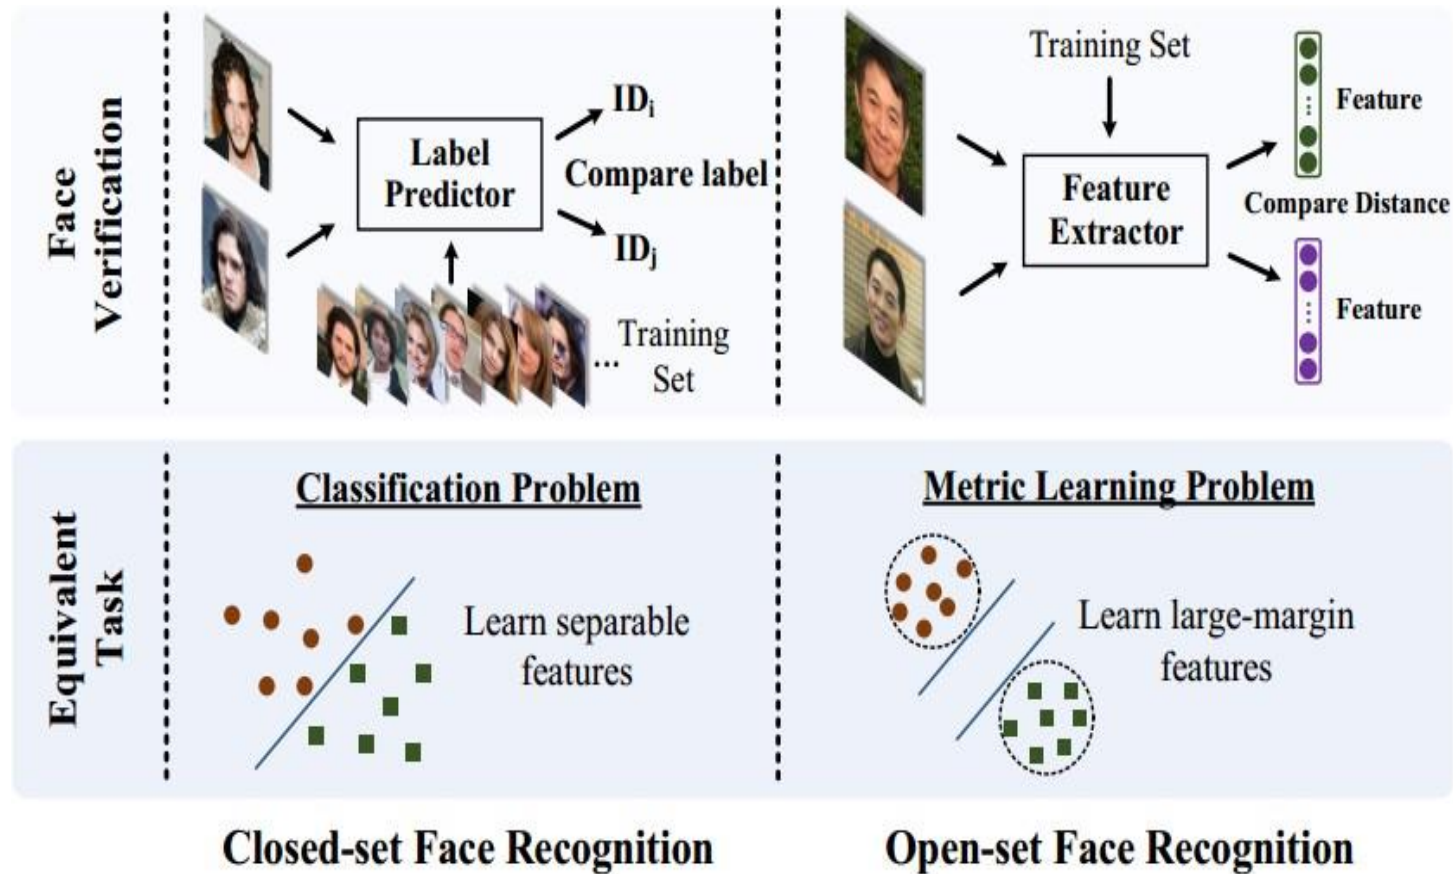
Motivation

Closed-set and Open-set Face Recognition



Motivation

Closed-set and Open-set Face Recognition



Softmax Loss

- Given input feature \mathbf{x}_i with the label y_i , the softmax loss function is:

$$\mathcal{L} = \frac{1}{N} \sum_i L_i = \frac{1}{N} \sum_i -\log \frac{e^{f_{y_i}}}{\sum_j e^{f_j}}$$

- f_j denotes the j -th element of the vector of class scores f
- N is the number of training data

$$f_{y_i} = \mathbf{w}_{y_i}^T \mathbf{x}_i = \|\mathbf{w}_{y_i}\| \|\mathbf{x}_i\| \cos(\theta_{y_i})$$

$$\mathcal{L}_i = -\log \left(\frac{e^{\|\mathbf{w}_{y_i}\| \|\mathbf{x}_i\| \cos(\theta_{y_i})}}{\sum_j e^{\|\mathbf{w}_j\| \|\mathbf{x}_i\| \cos(\theta_j)}} \right)$$

- θ_j ($0 \leq \theta_j \leq \pi$) is the angle between the vector \mathbf{w}_j and \mathbf{x}_i

Large-Margin Softmax Loss

- Consider the binary classification and a sample \mathbf{x} from class 1
- Original softmax

$$\|\mathbf{w}_1\| \|\mathbf{x}\| \cos(\theta_1) > \|\mathbf{w}_2\| \|\mathbf{x}\| \cos(\theta_2)$$

- Large-Margin softmax

$$\|\mathbf{w}_1\| \|\mathbf{x}\| \cos(m\theta_1) > \|\mathbf{w}_2\| \|\mathbf{x}\| \cos(\theta_2) \quad (0 \leq \theta_1 \leq \frac{\pi}{m})$$

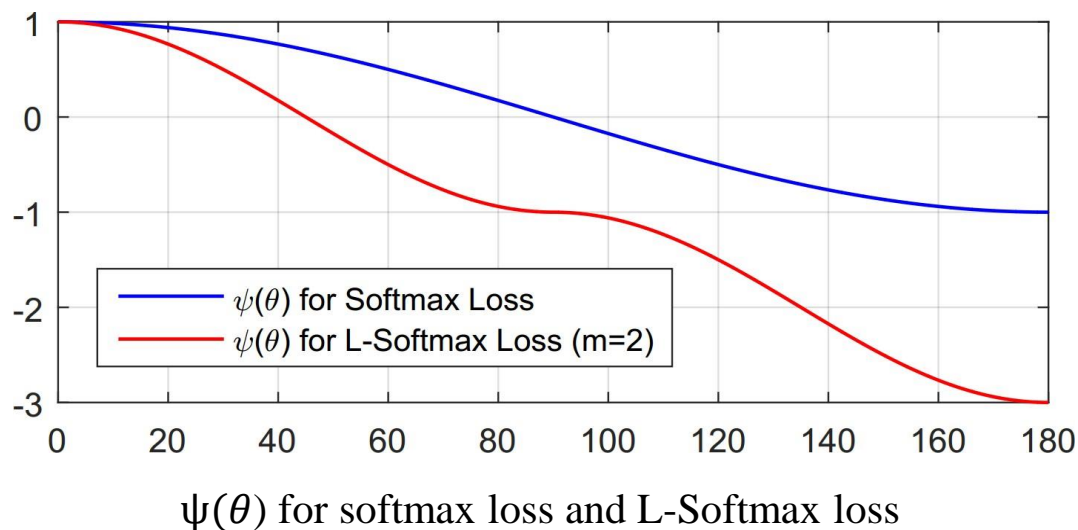
Large-Margin Softmax Loss

■ Large-Margin Softmax Loss:

$$L_i = -\log \left(\frac{e^{\|\mathbf{w}_{y_i}\| \|\mathbf{x}_i\| \psi(\theta_{y_i})}}{e^{\|\mathbf{w}_{y_i}\| \|\mathbf{x}_i\| \psi(\theta_{y_i})} + \sum_{j \neq y_i} e^{\|\mathbf{w}_{y_i}\| \|\mathbf{x}_i\| \cos(\theta_j)}} \right)$$

$$\psi(\theta) = \begin{cases} \cos(m\theta), & 0 \leq \theta \leq \frac{\pi}{m} \\ \mathcal{D}(\theta), & \frac{\pi}{m} \leq \theta \leq \pi \end{cases}$$

Large-Margin Softmax Loss



■ Construct a specific $\psi(\theta)$:

$$\psi(\theta) = (-1)^k \cos(m\theta) - 2k, \theta \in \left[\frac{k\pi}{m}, \frac{(k+1)\pi}{m} \right]$$

where $k \in [0, m-1]$ and k is an integer

Large-Margin Softmax Loss

- Replace $\cos(\theta_j)$ with

$$\frac{\mathbf{w}_j^T \mathbf{x}_i}{\|\mathbf{w}_j\| \|\mathbf{x}_i\|}$$

- Replace $\cos(m\theta_{y_i})$ with

$$\begin{aligned} \cos(m\theta_{y_i}) = & C_m^0 \cos^m(\theta_{y_i}) - C_m^2 \cos^{m-2}(\theta_{y_i}) \left(1 - \cos^2(\theta_{y_i})\right) \\ & + C_m^4 \cos^{m-4}(\theta_{y_i}) \left(1 - \cos^2(\theta_{y_i})\right)^2 + \dots \\ & (-1)^n C_m^{2n} \cos^{m-2n}(\theta_{y_i}) \left(1 - \cos^2(\theta_{y_i})\right)^n + \dots \end{aligned}$$

Large-Margin Softmax Loss

■ So we can get:

$$\begin{aligned} f_{y_i} &= (-1)^k \cdot \|\mathbf{w}_{y_i}\| \|\mathbf{x}_i\| \cos(m\theta_i) - 2k \cdot \|\mathbf{w}_{y_i}\| \|\mathbf{x}_i\| \\ &= (-1)^k \cdot \|\mathbf{w}_{y_i}\| \|\mathbf{x}_i\| \\ &\quad \cdot \left(C_m^0 \left(\frac{\mathbf{w}_{y_i}^T \mathbf{x}_i}{\|\mathbf{w}_{y_i}\| \|\mathbf{x}_i\|} \right)^m - C_m^2 \left(\frac{\mathbf{w}_{y_i}^T \mathbf{x}_i}{\|\mathbf{w}_{y_i}\| \|\mathbf{x}_i\|} \right)^{m-2} \left(1 - \left(\frac{\mathbf{w}_{y_i}^T \mathbf{x}_i}{\|\mathbf{w}_{y_i}\| \|\mathbf{x}_i\|} \right)^2 \right) + \dots \right) \\ &\quad - 2k \cdot \|\mathbf{w}_{y_i}\| \|\mathbf{x}_i\| \end{aligned}$$

where $\frac{\mathbf{w}_{y_i}^T \mathbf{x}_i}{\|\mathbf{w}_{y_i}\| \|\mathbf{x}_i\|} \in \left[\cos\left(\frac{k\pi}{m}\right), \cos\left(\frac{(k+1)\pi}{m}\right) \right]$ and k is an integer that to $[0, m-1]$.

Large-Margin Softmax Loss Optimization

$$\begin{aligned}\frac{\partial f_{y_i}}{\partial \mathbf{x}_i} = & (-1)^k \cdot (C_m^0 \left(\frac{m(\mathbf{w}_{y_i}^T \mathbf{x}_i)^{m-1} \mathbf{w}_{y_i}}{(\|\mathbf{w}_{y_i}\| \|\mathbf{x}_i\|)^{m-1}} \right) - \\ & C_m^0 \left(\frac{(m-1)(\mathbf{w}_{y_i}^T \mathbf{x}_i)^m \mathbf{x}_i}{\|\mathbf{w}_{y_i}\|^{m-1} \|\mathbf{x}_i\|^{m+1}} \right) - C_m^2 \left(\frac{(m-2)(\mathbf{w}_{y_i}^T \mathbf{x}_i)^{m-3} \mathbf{w}_{y_i}}{(\|\mathbf{w}_{y_i}\| \|\mathbf{x}_i\|)^{m-3}} \right) \\ & + C_m^2 \left(\frac{(m-3)(\mathbf{w}_{y_i}^T \mathbf{x}_i)^{m-2} \mathbf{x}_i}{\|\mathbf{w}_{y_i}\|^{m-3} \|\mathbf{x}_i\|^{m-1}} \right) + C_m^2 \left(\frac{m(\mathbf{w}_{y_i}^T \mathbf{x}_i)^{m-1} \mathbf{w}_{y_i}}{(\|\mathbf{w}_{y_i}\| \|\mathbf{x}_i\|)^{m-1}} \right) \\ & - C_m^2 \left(\frac{(m-1)(\mathbf{w}_{y_i}^T \mathbf{x}_i)^m \mathbf{x}_i}{\|\mathbf{w}_{y_i}\|^{m-1} \|\mathbf{x}_i\|^{m+1}} \right) + \dots) - 2k \cdot \frac{\|\mathbf{w}_{y_i}\| \mathbf{x}_i}{\|\mathbf{x}_i\|}\end{aligned}$$

Large-Margin Softmax Loss

Optimization

$$\begin{aligned}
 \frac{\partial f_{y_i}}{\partial \mathbf{w}_{y_i}} = & (-1)^k \cdot (C_m^0 \left(\frac{m(\mathbf{w}_{y_i}^T \mathbf{x}_i)^{m-1} \mathbf{x}_i}{(\|\mathbf{w}_{y_i}\| \|\mathbf{x}_i\|)^{m-1}} \right) \\
 & - C_m^0 \left(\frac{(m-1)(\mathbf{w}_{y_i}^T \mathbf{x}_i)^m \mathbf{w}_{y_i}}{\|\mathbf{w}_{y_i}\|^{m+1} \|\mathbf{x}_i\|^{m-1}} \right) \\
 & - C_m^2 \left(\frac{(m-2)(\mathbf{w}_{y_i}^T \mathbf{x}_i)^{m-3} \mathbf{x}_i}{(\|\mathbf{w}_{y_i}\| \|\mathbf{x}_i\|)^{m-3}} \right) \\
 & + C_m^2 \left(\frac{(m-3)(\mathbf{w}_{y_i}^T \mathbf{x}_i)^{m-2} \mathbf{w}_{y_i}}{\|\mathbf{w}_{y_i}\|^{m-1} \|\mathbf{x}_i\|^{m-3}} \right) \\
 & + C_m^2 \left(\frac{m(\mathbf{w}_{y_i}^T \mathbf{x}_i)^{m-1} \mathbf{x}_i}{(\|\mathbf{w}_{y_i}\| \|\mathbf{x}_i\|)^{m-1}} \right) - C_m^2 \left(\frac{(m-1)(\mathbf{w}_{y_i}^T \mathbf{x}_i)^m \mathbf{w}_{y_i}}{\|\mathbf{w}_{y_i}\|^{m+1} \|\mathbf{x}_i\|^{m-1}} \right) \\
 & + \dots) - 2k \cdot \frac{\|\mathbf{x}_i\| \mathbf{w}_{y_i}}{\|\mathbf{w}_{y_i}\|}
 \end{aligned}$$

Geometric Interpretation




Figure: Example of Geometric Interpretation when $\|\mathbf{w}_1\| = \|\mathbf{w}_2\|$

Geometric Interpretation

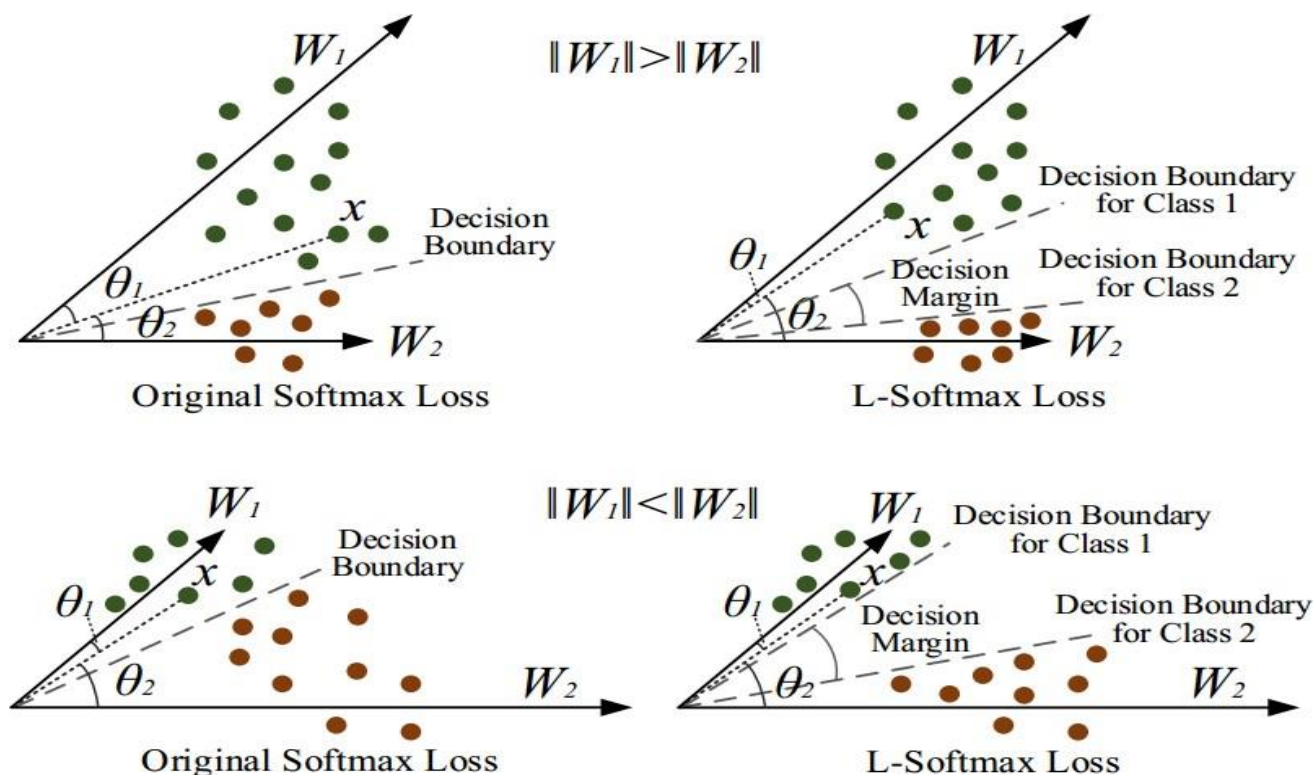


Figure: Example of Geometric Interpretation when $\|\mathbf{w}_1\| > \|\mathbf{w}_2\|$ and $\|\mathbf{w}_1\| < \|\mathbf{w}_2\|$

The variants of the softmax loss

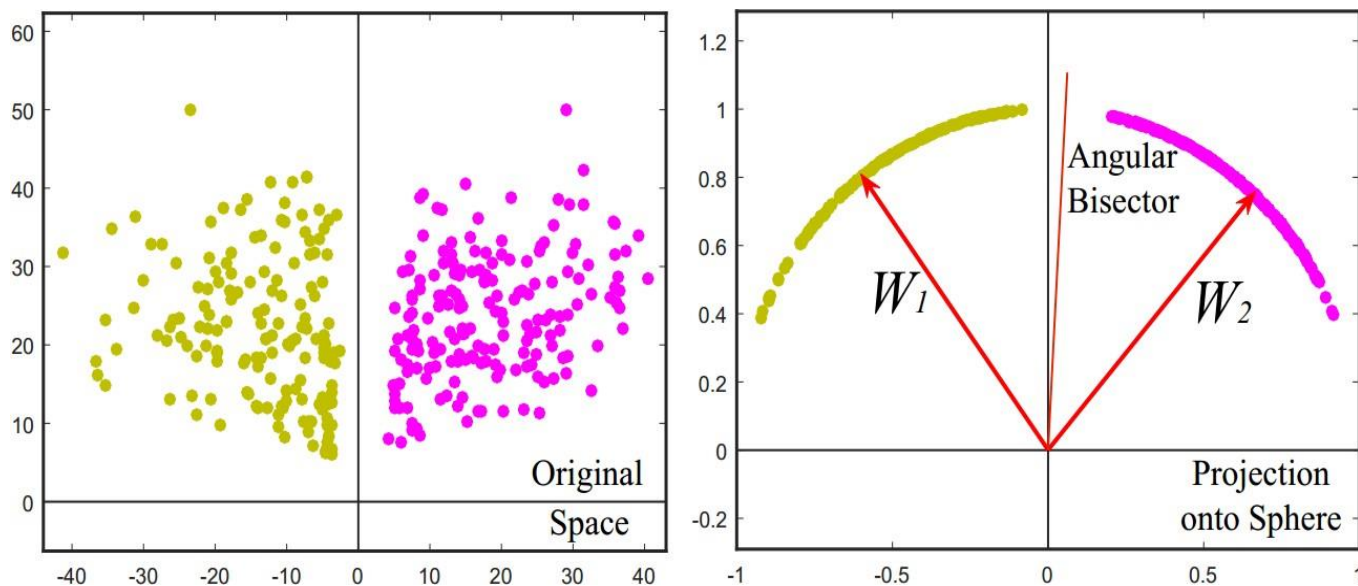
- Large-Margin Softmax Loss
- **Angular Softmax Loss (A-Softmax Loss)**

Modified Softmax Loss Function

- Normalize $\|\mathbf{w}_j\| = 1, \forall j$ in each iteration

$$\mathcal{L}_{modified} = \frac{1}{N} \sum_i -\log\left(\frac{e^{\|\mathbf{x}_i\| \cos(\theta_{y_i, i})}}{\sum_j e^{\|\mathbf{x}_i\| \cos(\theta_{j, i})}}\right)$$

Modified Softmax Loss Function



- Learn a 2-D features on subset of CASIA face dataset

A-Softmax Loss

- Consider the binary classification and a sample x from class 1
- Modified softmax loss need

$$\|x\| \cos(\theta_1) > \|x\| \cos(\theta_2)$$

- A-Softmax loss need

$$\|x\| \cos(m\theta_1) > \|x\| \cos(\theta_2) \quad (0 \leq \theta_1 \leq \frac{\pi}{m})$$

A-Softmax Loss


$$L_{ang} = \frac{1}{N} \sum_i -\log\left(\frac{e^{\|\mathbf{x}_i\| \cos(m\theta_{y_i,i})}}{e^{\|\mathbf{x}_i\| \cos(m\theta_{y_i,i})} + \sum_{j \neq y_i} e^{\|\mathbf{x}_i\| \cos(\theta_{j,i})}}\right)$$

where $\theta_{y_i,i}$, i has to be in the range of $[0, \frac{\pi}{m}]$

$$L_{ang} = \frac{1}{N} \sum_i -\log\left(\frac{e^{\|\mathbf{x}_i\| \psi(\theta_{y_i,i})}}{e^{\|\mathbf{x}_i\| \psi(\theta_{y_i,i})} + \sum_{j \neq y_i} e^{\|\mathbf{x}_i\| \cos(\theta_{j,i})}}\right)$$

$$\psi(\theta) = \begin{cases} \cos(m\theta), & 0 \leq \theta \leq \frac{\pi}{m} \\ \mathcal{D}(\theta), & \frac{\pi}{m} \leq \theta \leq \pi \end{cases}$$

A-Softmax Loss




■ Construct a specific $\psi(\theta)$:

$$\psi(\theta) = (-1)^k \cos(m\theta) - 2k, \theta \in \left[\frac{k\pi}{m}, \frac{(k+1)\pi}{m}\right]$$


where $k \in [0, m-1]$ and k is an integer

A-Softmax Loss

Geometric Interpretation



Modified Softmax Loss



A-Softmax Loss

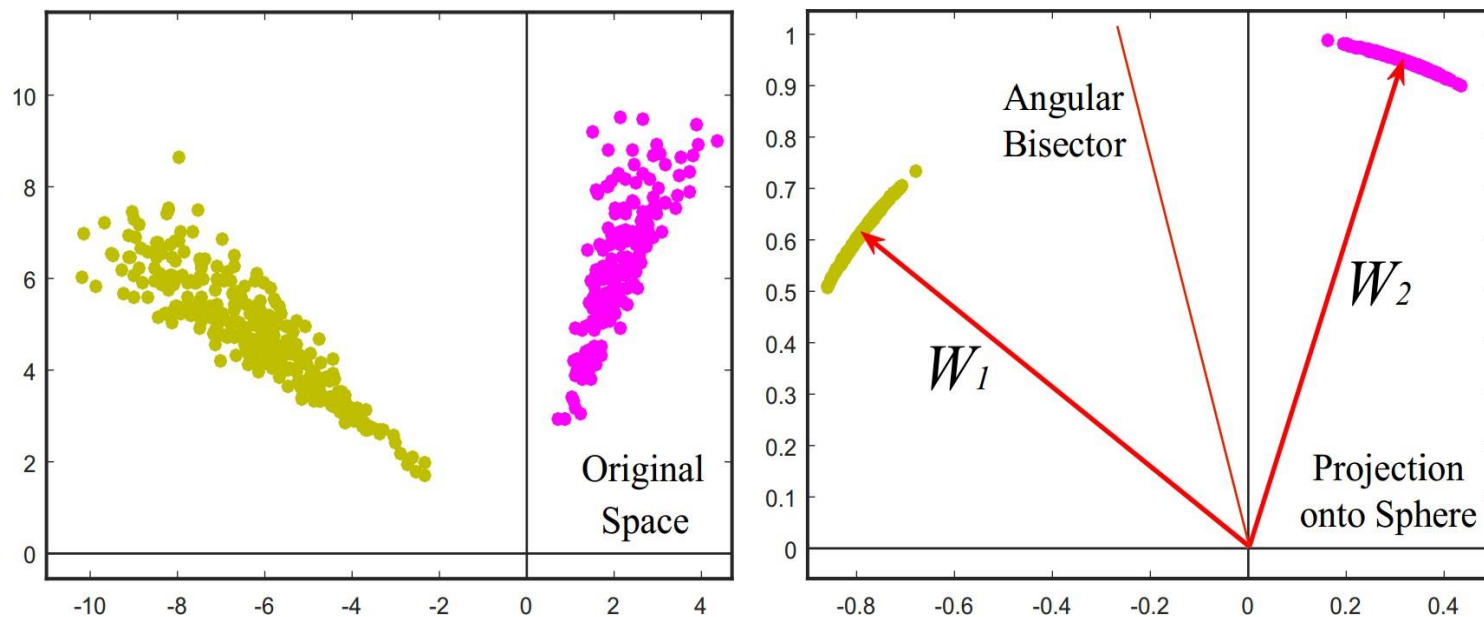
A-Softmax Loss

Decision Boundary

Loss Function	Decision Boundary
Softmax Loss	$(\mathbf{W}_1 - \mathbf{W}_2)\mathbf{x} + b_1 - b_2 = 0$
Modified Softmax Loss	$\ \mathbf{x}\ (\cos \theta_1 - \cos \theta_2) = 0$
A-Softmax Loss	$\ \mathbf{x}\ (\cos m\theta_1 - \cos \theta_2) = 0$ for class 1 $\ \mathbf{x}\ (\cos \theta_1 - \cos m\theta_2) = 0$ for class 2



- θ_i is the angle between \mathbf{w}_i and \mathbf{x}

A-Softmax Loss





- Learn 2-D features on a subset of CASIA face dataset



Hypersphere Interpretation



Euclidean Margin Loss



Modified Softmax Loss



A-Softmax Loss ($m \geq 2$)

References

- [1] Liu W, Wen Y, Yu Z, et al. Large-margin softmax loss for convolutional neural networks[C]//ICML. 2016, 2(3): 7.
- [2] Liu W, Wen Y, Yu Z, et al. Sphreface: Deep hypersphere embedding for face recognition[C] //Proceedings of the IEEE conference on computer vision and pattern recognition. 2017: 212-220.

Thank You