

Underfitting, Overfitting and Cross-Validation

Prof. Mingkui Tan

SCUT Machine Intelligence Laboratory (SMIL)



Contents

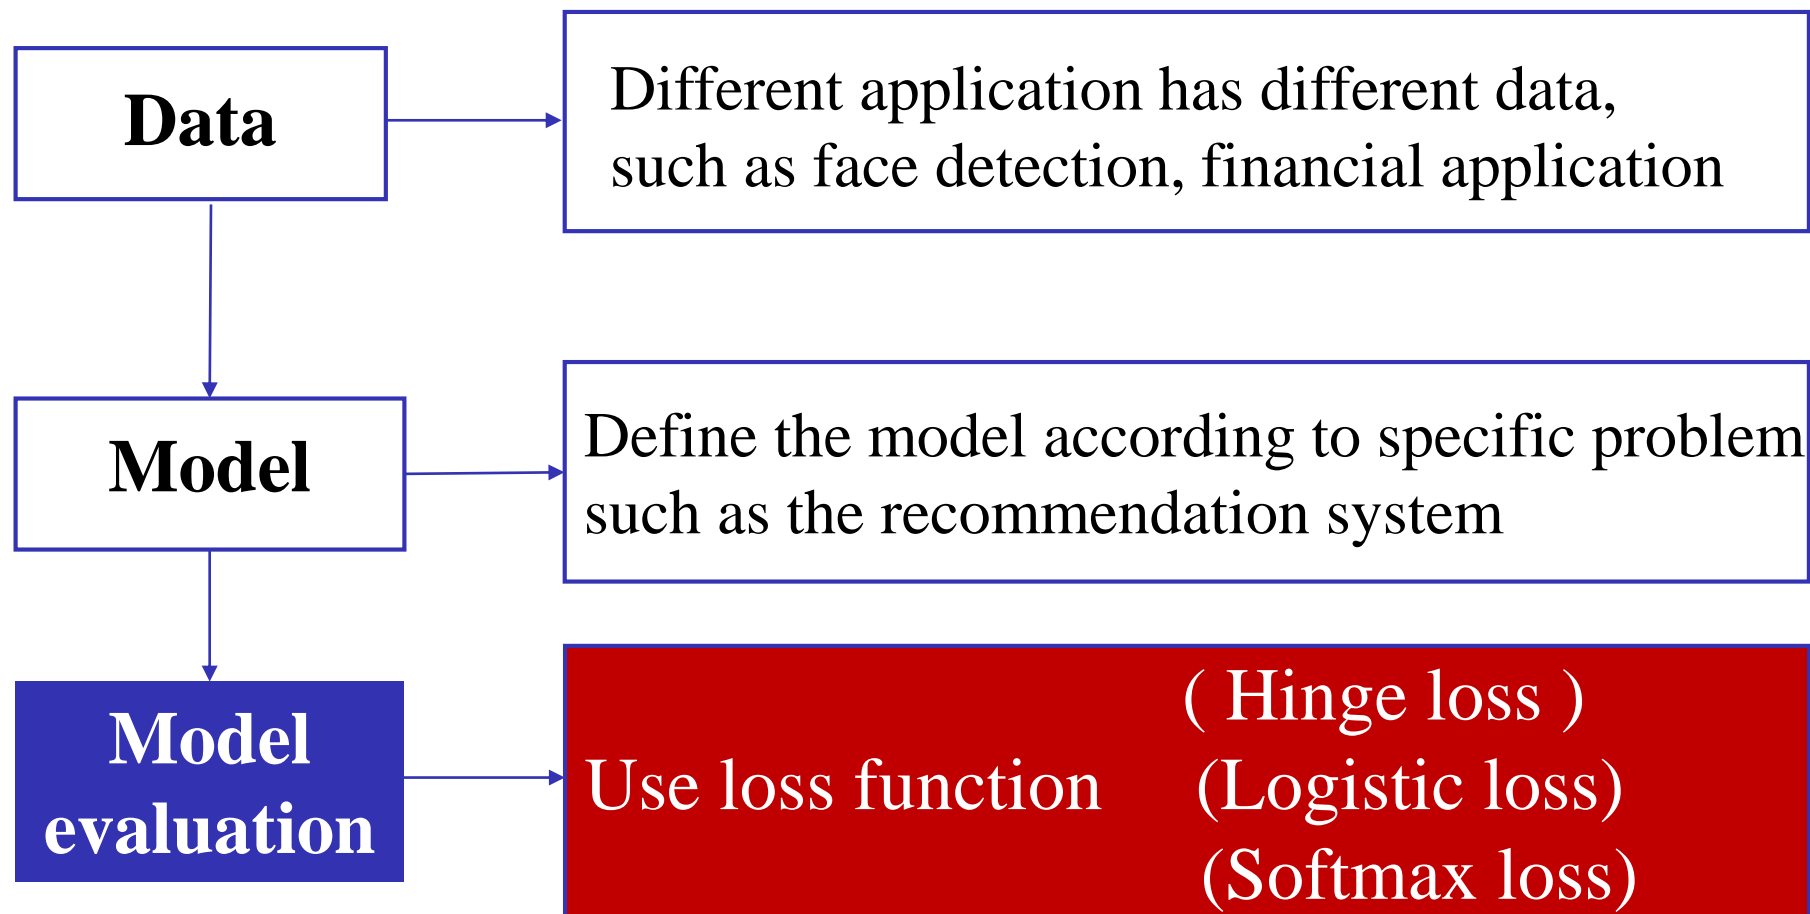
1 Training, Testing and Validation Split

2 Underfitting and Overfitting

3 Bias-Variance Trade-off

3 Cross-Validation

Main Elements of Machine Learning



Given a data set D , how to evaluate the performance of a learned model?

Data Notation

Data:

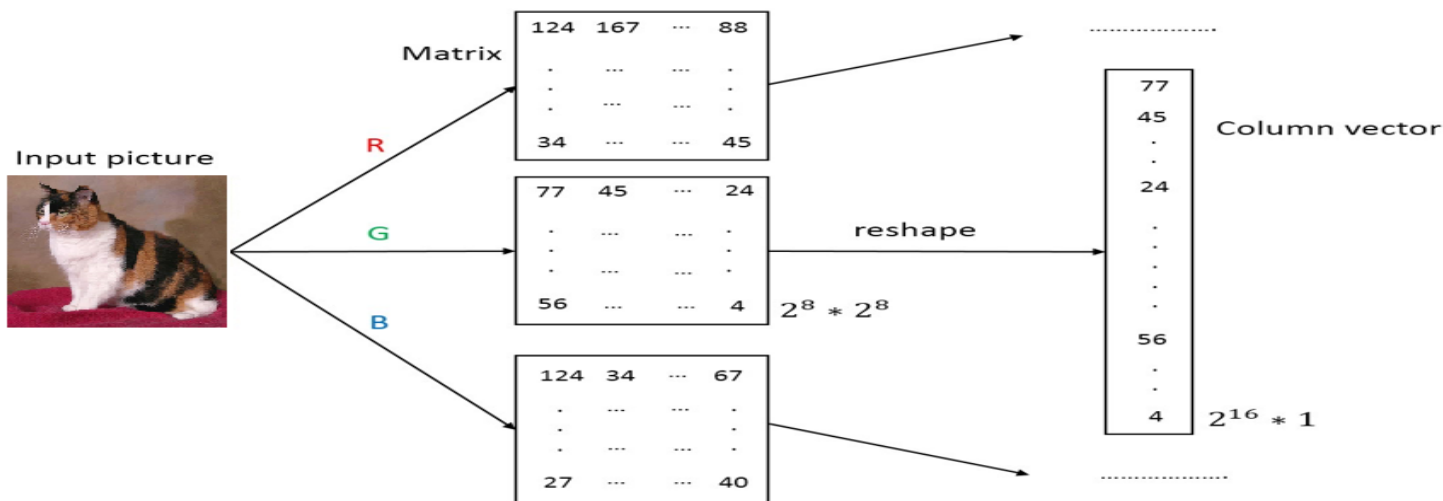
$$\mathcal{D} = \{(\mathbf{x}_i, y_i)\}_{i=1}^n$$

\mathbf{x} is the input, which is usually presented as a **column vector**

y is the output, for example, a person's name


n is the number of samples

For example, \mathbf{x} can be a picture stored as a matrix:



Regression Example

Example: small error variance



Regression

Loss:

- Absolute value loss:

$$l(\hat{y}_i, y_i) = |\hat{y}_i - y_i|$$

- Least squares loss:

$$l(\hat{y}_i, y_i) = \frac{1}{2} (\hat{y}_i - y_i)^2$$

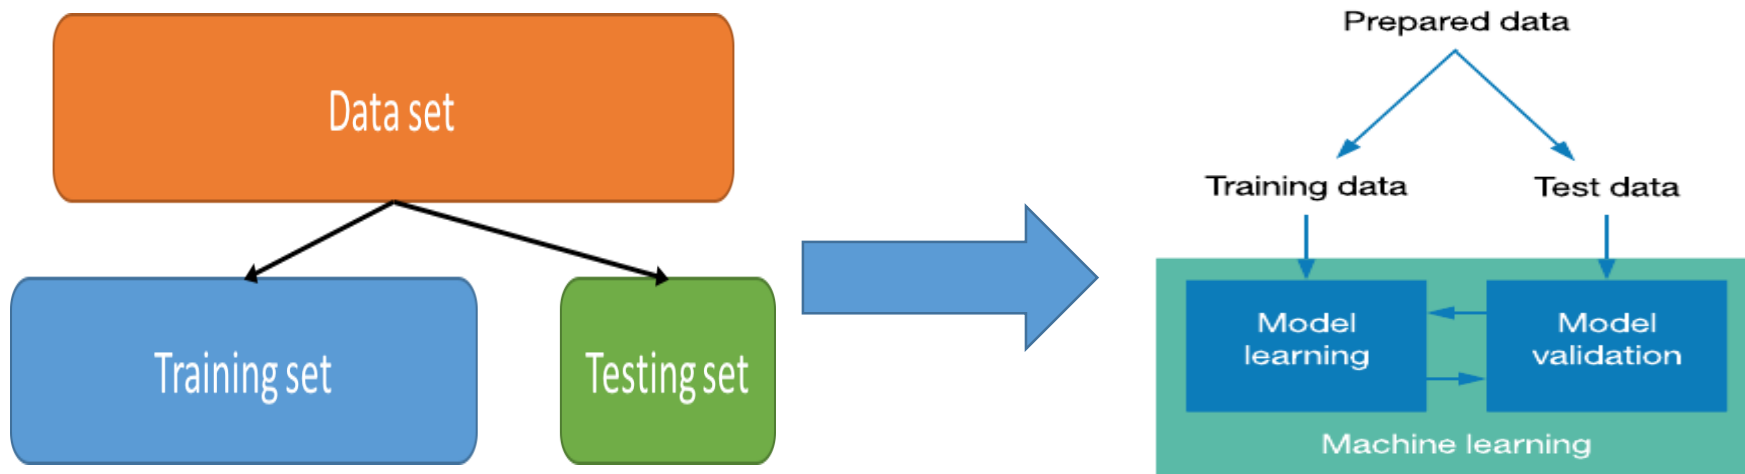
Given a data set D , how to evaluate the performance of a learned model?

Training-Testing-Validation Data Split

- Simple idea: Split data into two sets
 - **Training set:** Data used to fit the model
 - **Testing set:** Data used to evaluate the model



Data Split



■ Training set

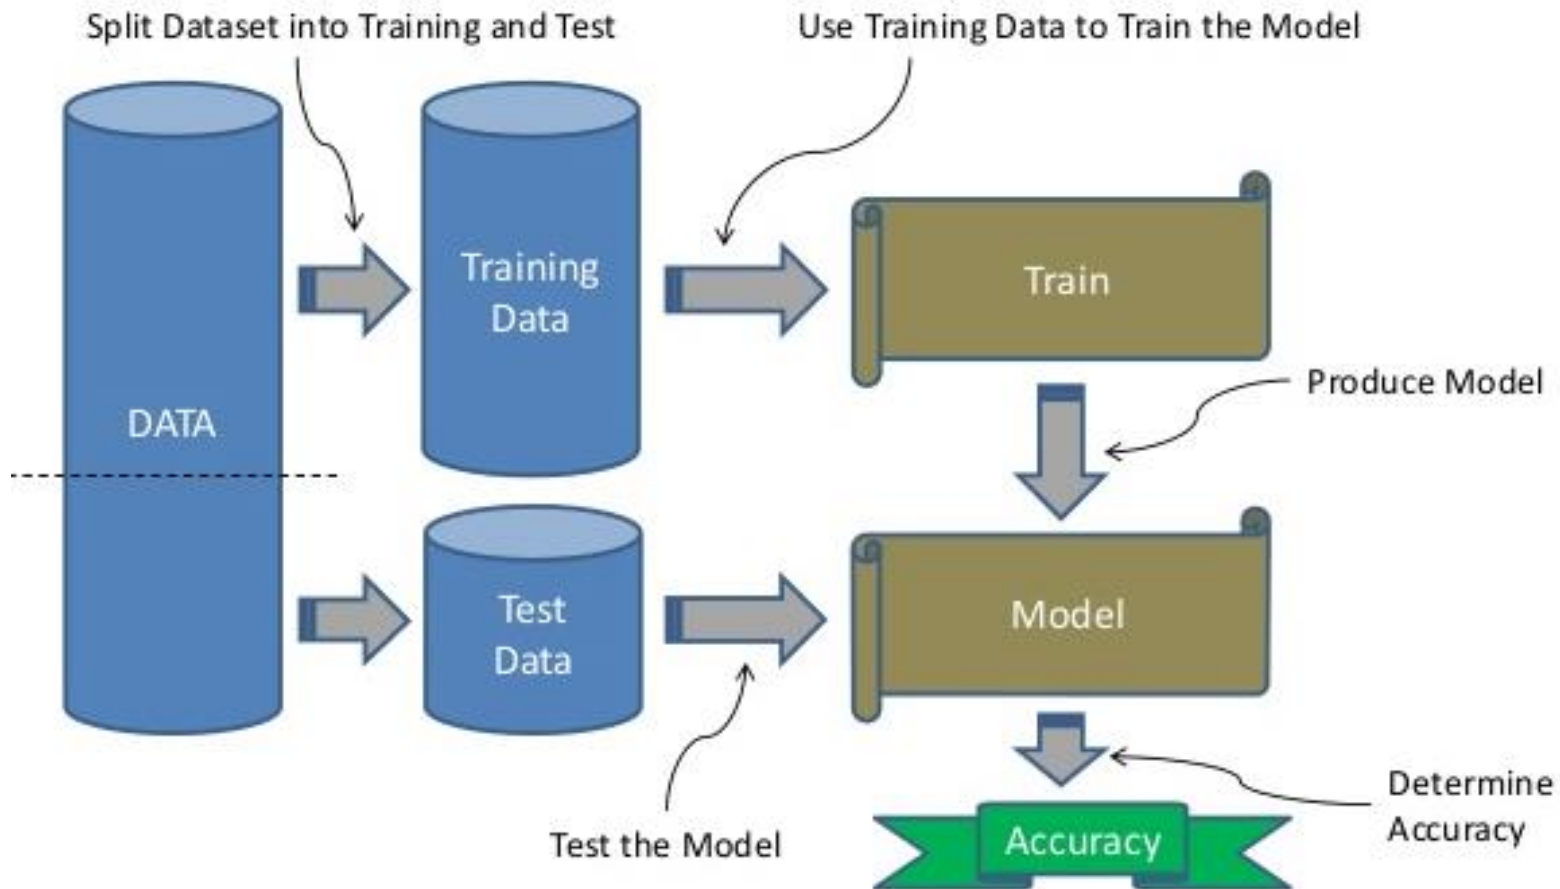
■ $\approx 70\%$ of data, $D_{train} = \{x^{(i)}, y^{(i)}\}$, total n_{train} examples

■ Testing set

■ $\approx 30\%$ of data, $D_{test} = \{x_{test}^{(i)}, y_{test}^{(i)}\}$, total n_{test} examples

■ Choose examples randomly for training/testing split

Data Split for Training and Testing



Train-test split example

```
1 # train-test split evaluation random forest on the housing dataset
2 from pandas import read_csv
3 from sklearn.model_selection import train_test_split
4 from sklearn.ensemble import RandomForestRegressor
5 from sklearn.metrics import mean_absolute_error
6 # load dataset
7 url = 'https://raw.githubusercontent.com/jbrownlee/Datasets/master/housing.csv'
8 dataframe = read_csv(url, header=None)
9 data = dataframe.values
10 # split into inputs and outputs
11 X, y = data[:, :-1], data[:, -1]
12 print(X.shape, y.shape)
13 # split into train test sets
14 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, random_state=1)
15 print(X_train.shape, X_test.shape, y_train.shape, y_test.shape)
16 # fit the model
17 model = RandomForestRegressor(random_state=1)
18 model.fit(X_train, y_train)
19 # make predictions
20 yhat = model.predict(X_test)
21 # evaluate predictions
22 mae = mean_absolute_error(y_test, yhat)
23 print('MAE: %.3f' % mae)
```

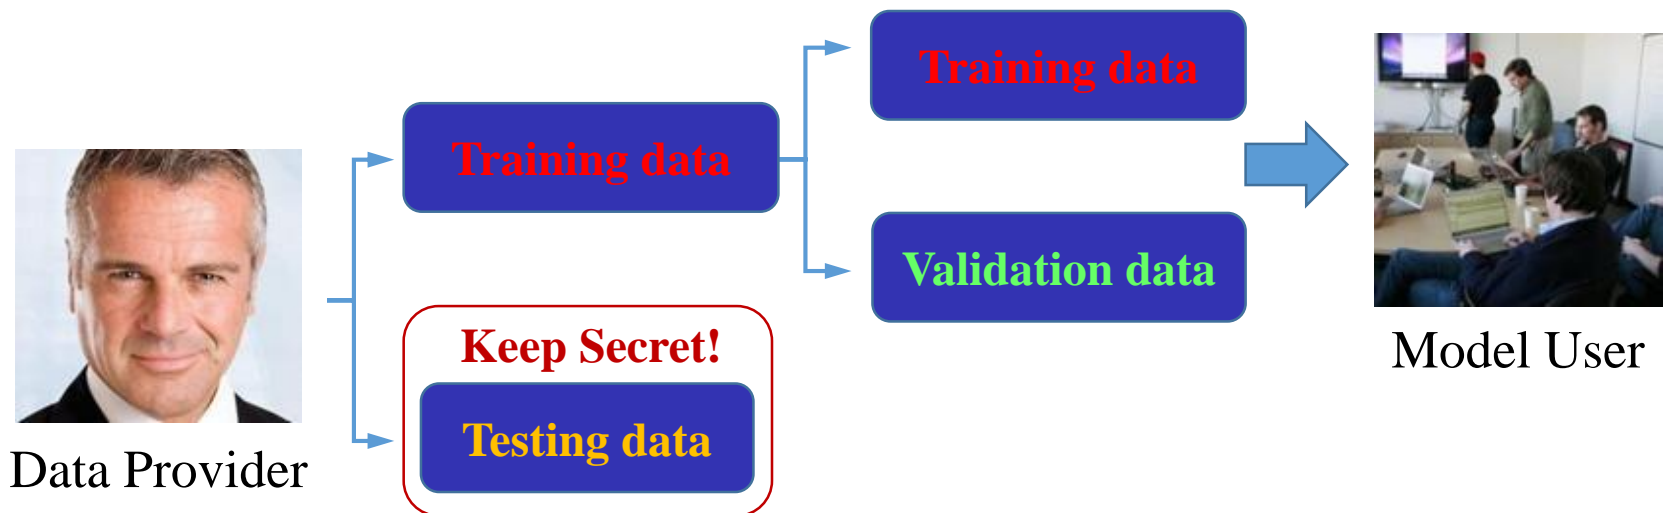
Source available:

<https://machinelearningmastery.com/train-test-split-for-evaluating-machine-learning-algorithms/>

Training-Testing-Validation Data Split

■ Split data into three sets


- We still have training and testing sets
- But additionally, we have a **validation set** to test the performance of our model depending on the parameter



Why We Need Validation Set?

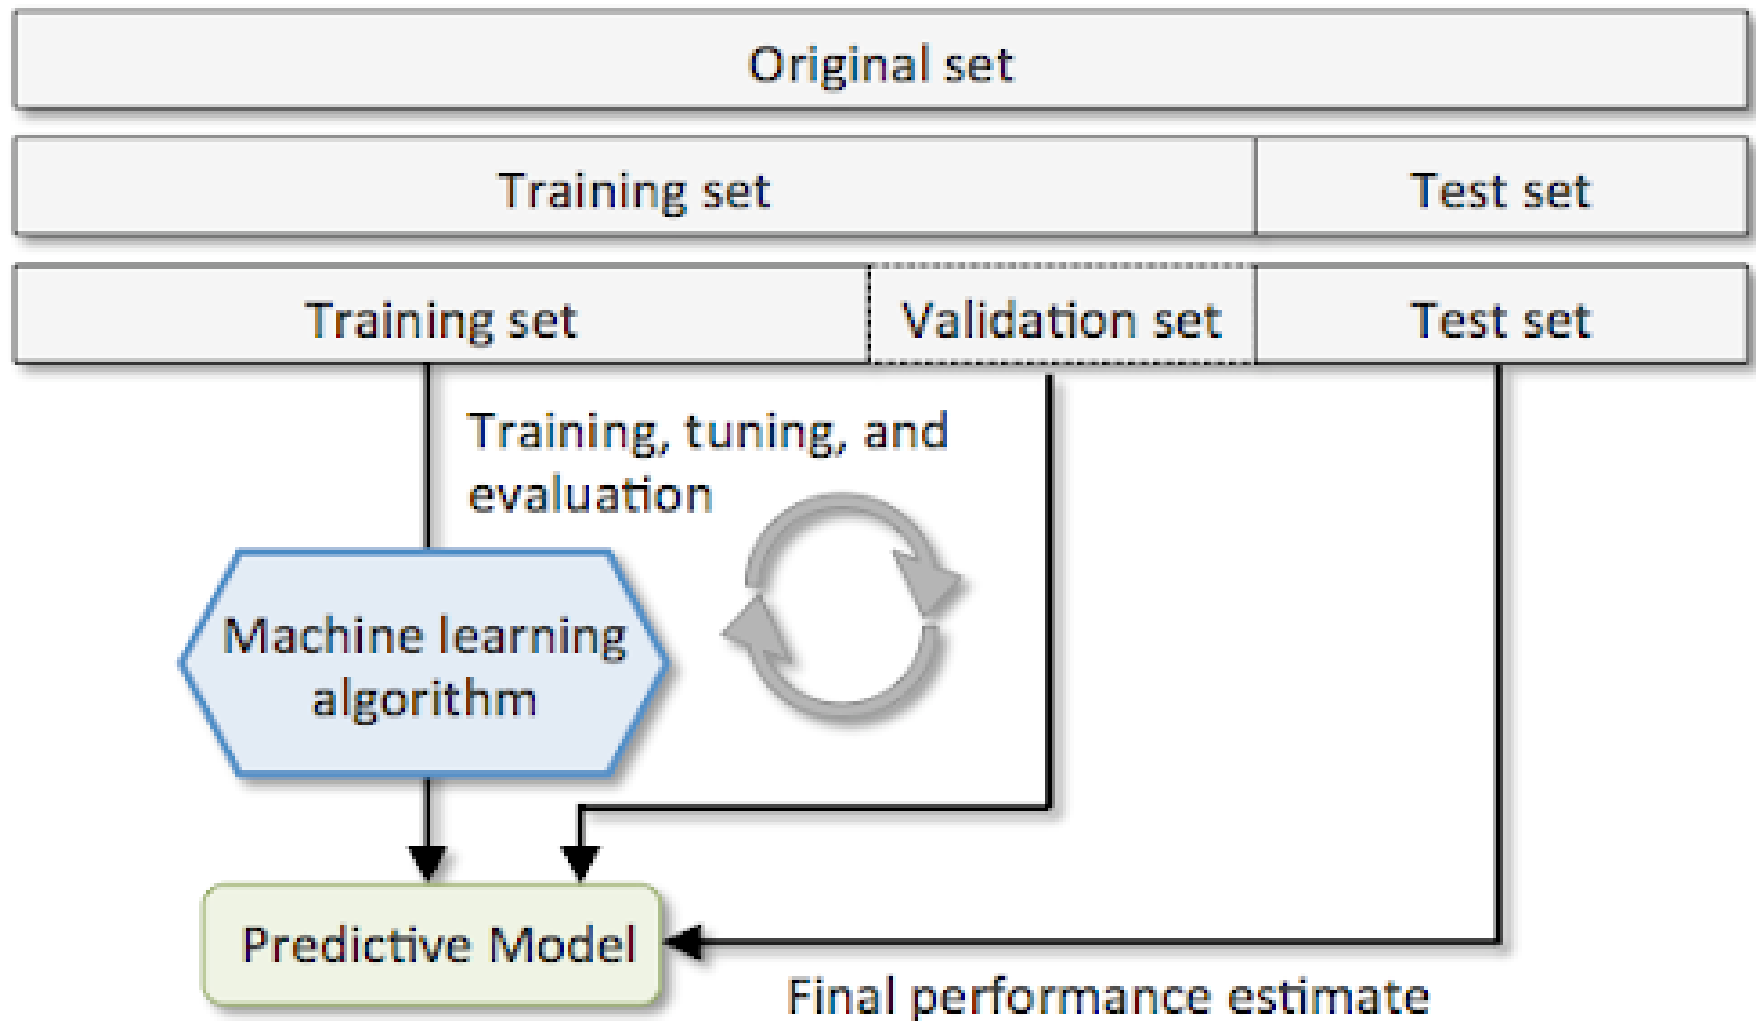


Get Test set error



CV Loop
Tune hyperparameters

Why We Need Validation Set?



Why We Need Validation Set?

■ Business Reasons

- Need to choose the best model
- Measure accuracy/power of the selected model
- Better to measure ROI of the modeling project

■ Statistical Reasons

- Model building techniques are inherently designed to minimize “loss” or “bias”
- To an extent, a model will always fit “noise” as well as “signal”
- If you just fit a bunch of models on a given dataset and choose the “best” one, it will likely be overly “optimistic”

Contents


1 Underfitting and Overfitting

2 Bias-Variance Trade-off

3 Cross-Validation


General Fitting Scheme

- Model is fitting and can capture the underlying trend of the data




Underfitting

- Model cannot capture the underlying trend of the data




Overfitting

- The model is too complex to capture the true trend
- The model seeks to fit the noise or outlier of the data




Underfitting vs Overfitting




$$g(\theta_0 + \theta_1 x_1 + \theta_2 x_2)$$

(g = sigmoid function)



$$g(\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_1^2 + \theta_4 x_2^2)$$



$$g(\theta_0 + \theta_1 x_1 + \theta_2 x_1^2 + \theta_3 x_1^2 x_2 + \theta_4 x_1^2 x_2^2 + \theta_5 x_1^2 x_2^3 + \theta_6 x_1^3 x_2 + \dots)$$

■ Comprehend from Taylor expansion

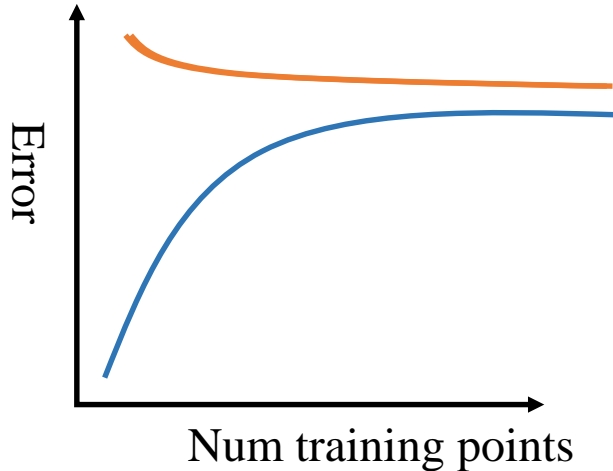
- The more terms, the more complex, the more power

$$f(x) = f(x_0) + \nabla f(x_0)^T (x - x_0) + \frac{1}{2!} (x - x_0)^T \nabla^2 f(x_0) (x - x_0) + \dots$$


Signs of Underfitting and Overfitting

- How to judge underfitting or overfitting?
 - Underfitting: If the training set's error is relatively large and the generalization error is large
 - Need to increase capacity (complexity of models)
 - Overfitting: If the training set's error is relatively small and the generalization error is large
 - Need to decrease capacity (complexity of models)
 - Or increase training set

Signs of Underfitting and Overfitting



High Bias
(Underfitting)



High Variance
(Overfitting)

Contents

1 Underfitting and Overfitting


2 Bias-Variance Trade-off

3 Cross-Validation

Bias-Variance Trade-off

■ Complex model

- Too complex can diminish the model's accuracy on future data (called the Bias-Variance Trade-off)
- Low Bias
 - Model fits well on the training data
- High Variance
 - Model is more likely to make a wrong prediction



Contents

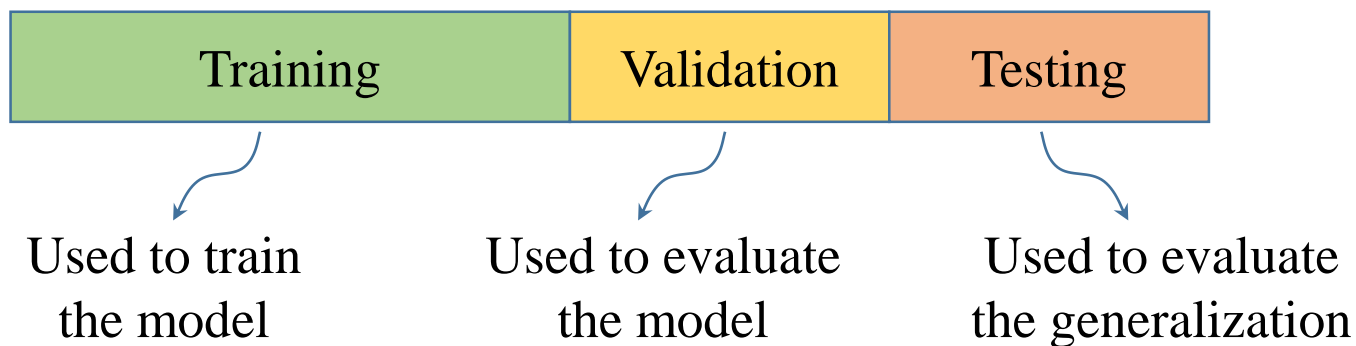
1 Underfitting and Overfitting

2 Bias-Variance Trade-off

3 Cross-Validation

Tuning Learning Parameter

- Use validation set for tuning hyper-parameters
- Use testing set only for final evaluation

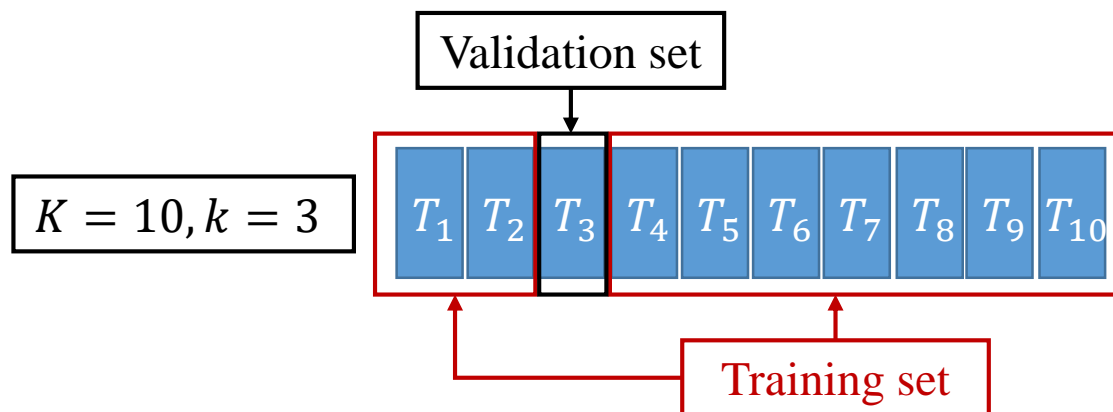


K-Fold Cross-validation

- If we want to reduce variability in the data
 - We can use multiple rounds of cross-validation using different partitions
 - And then, average the result overall rounds
- Given a data S sampled from the population D

K-Fold Cross-validation

- Split data S into K equal disjoint subsets (T_1, \dots, T_K)
- Perform the following steps for $k = 1, \dots, K$
 - Use $R_k = S - T_k$ as the training set
 - Build classifier C_k using R_k
 - Use T_k as the validation set, compute error $Err_k = error(C_k, T_k)$
- Let $Err^{ave} = \frac{1}{K} \sum_{k=1}^K Err_k$
 - This is the averaged error rate



Validation for Evaluation

■ Training error

$$J_{train}(\theta) = \frac{1}{2m} \sum cost(x^{(i)}, y^{(i)})$$

■ Validation error

$$J_{cv}(\theta) = \frac{1}{2m_{cv}} \sum cost(x_{cv}^{(i)}, y_{cv}^{(i)})$$

■ For model selection

- Obtain $\theta^{(1)}, \dots, \theta^{(d)}$ and select the best (lowest) $J_{cv}(\theta^{(i)})$

■ Testing error

$$J_{test}(\theta) = \frac{1}{2m_{test}} \sum cost(x_{test}^{(i)}, y_{test}^{(i)})$$

■ For final evaluation

- Evaluate generalization error $J_{test}(\theta^{(i)})$ on the testing set

Tuning Regularization Parameter λ

- Suppose we are fitting a model with high-order polynomial


$$h_{\theta}(x) = \theta_0 + \theta_1 x + \cdots + \theta_n x^n$$

- To prevent overfitting, we use regularization


$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \text{cost}(h_{\theta}(x_i), y_i) + \frac{\lambda}{2m} \sum_{j=1}^n \theta_j^2$$

Tuning Regularization Parameter λ


- If λ is too large, all θ are penalized and $\theta_1 \approx \theta_2 \approx \dots \approx 0$, $h_{\theta}(x) \approx \theta_0$
- If λ is intermediate, the model fits well
- If λ is too small, the model fits too well, i.e. overfitting



λ is large



λ is intermediate



λ is small

Tuning Regularization Parameter λ

- Choose good λ on the validation set
 - Choose a range of possible values for $\lambda(0.02, \dots, 0.24)$
 - That gives us 12 models with different parameters λ to check
 - For each λ_i :
 - Learn θ_i
 - Calculate $J_{cv}(\theta_i)$
 - Take λ_i with lowest $J_{cv}(\theta_i)$
 - Finally, we report the test error as $J_{test}(\theta_i)$

Tuning Learning Parameter

- Choosing λ with K-Fold Cross-validation
 - Split your data into training, validation, and testing set
 - For every possible value λ , estimate the error rate
 - Select λ with least average error rate Err^{ave}
 - Final evaluation of the testing set

Thank You